

Special Topics: Machine Learning (ML) for Networking

COL867

Holi, 2025

Network Telemetry

Tarun Mangla

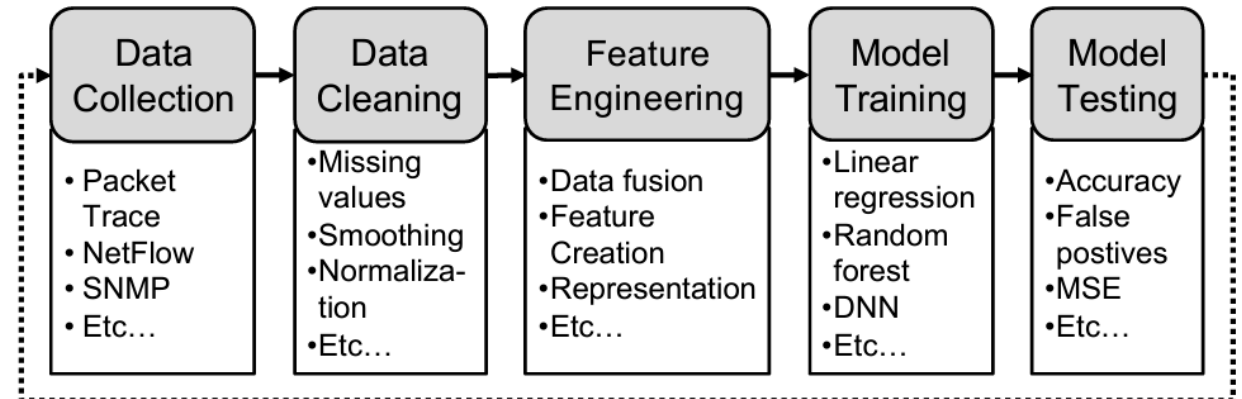
ML for Networks

Module 1: Case studies of specific network learning tasks

Module 2: Task-agnostic automatic ML pipelines for networks

Module 3: Beyond feature engineering and modeling

- • Network telemetry
- Robustness and explainability
- Synthetic data generation*
- Data imputation
- Formal verification
- ..




Network Telemetry

- ✓ • Network telemetry: Collection and analysis of network data
Fundamental to ML for networking
 - Given a network learning task, how would you collect network data?
 - Two approaches:
 - ○ Endogenous data collection
 - In controlled network conditions
- Challenge:
→ at scale & flexibility
- No noise
↳ Easy to label
- ↓
LIMITATION: ① May not represent the actual distribution

Significant storage & N/w overhead

Conventional Endogenous Network Telemetry

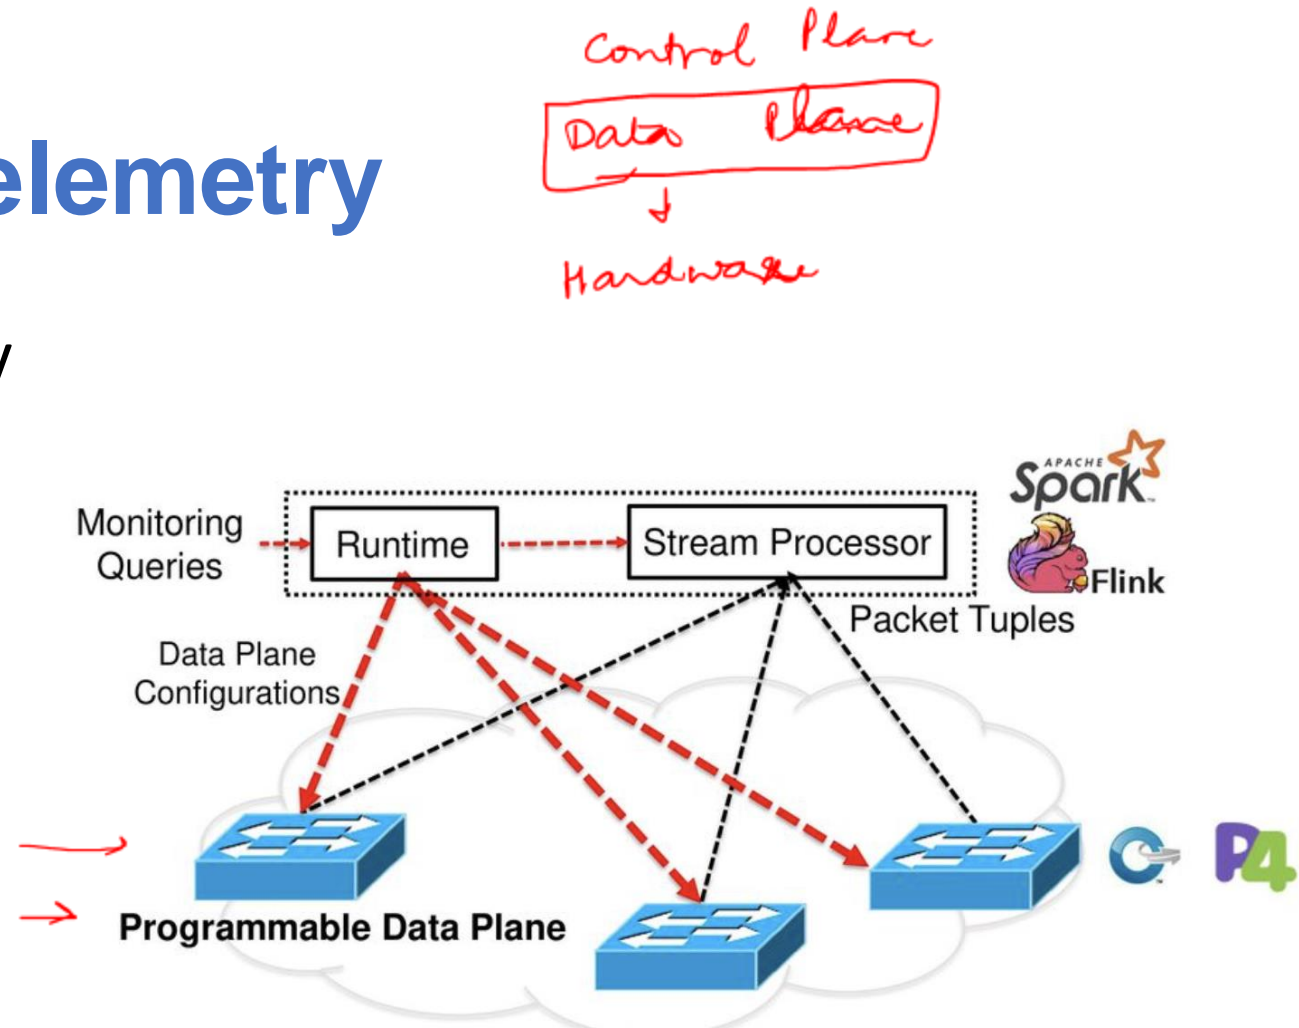
- Bottom-up data collection
 - Monitoring at routers and analysis at centralized servers
- Limitations:
 - • Flexibility vs scalability
- How to mitigate this cost?



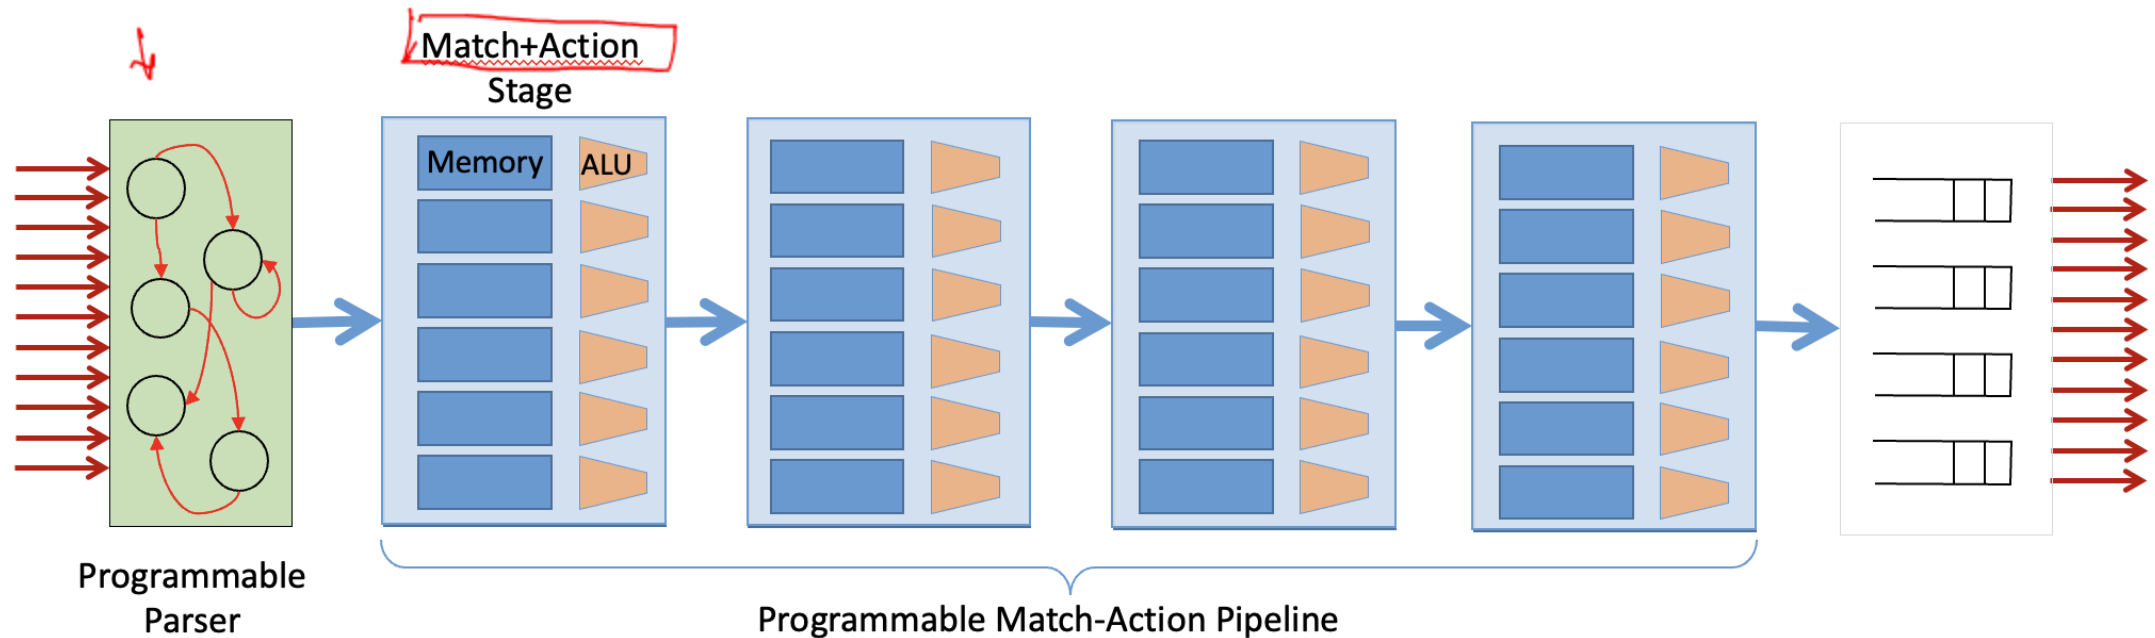
Top-down Network Telemetry

- Query-driven network telemetry
- Processing in the data plane
- But aren't routers **only** forwarding devices?

Not true anymore



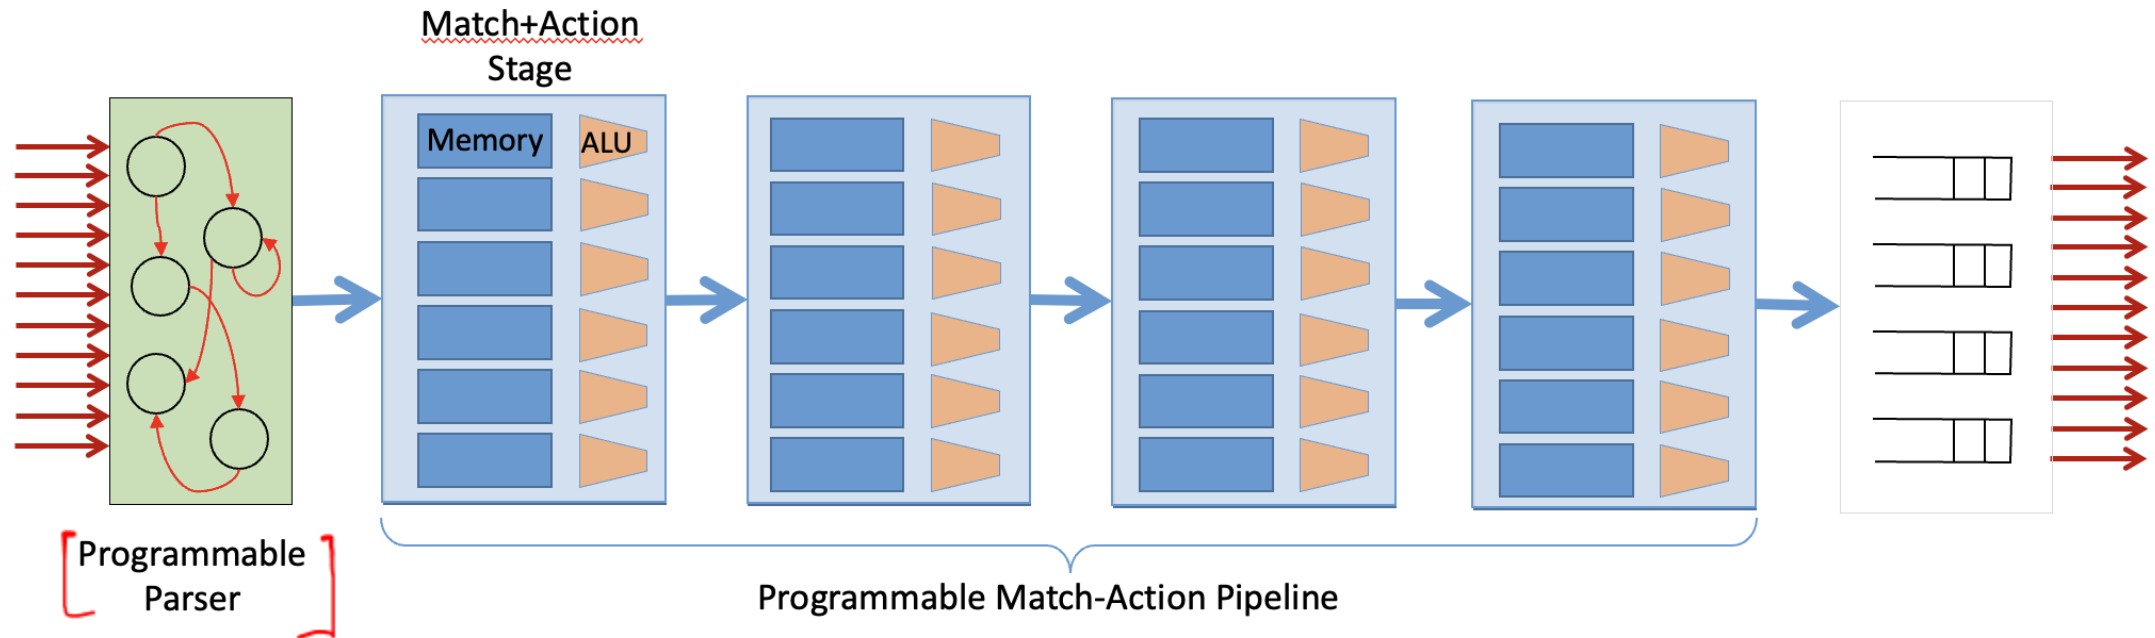
Programmable Data Plane Architecture



Key functionalities

- Programmable parser, extract specific fields of interest
- Matching, to group related packets into a single 'flow'
- Computation using simple ALUs (addition, subtraction, bitwise operations)
- Storage of information across successive packets using a small number of register arrays
- Communication with a software controller

Telemetry using Programmable Data Plane



How to calculate the number of packets per flow?

- ① Extract the 5-tuple
- ② 1 register for each 5-tuple

TCP: Close signal
FIN/RST

UDP:

Programmable Data Plane Constraints

- Memory is the biggest constraint
 - Packets forwarded without going through the main memory
 - Hence, measurements need to be performed in the register arrays
- ALUs available are simple
 - Can not perform complex operations
- Measurement is typically a second-class citizen
 - Gets limited budget compared to routing

How to calculate the number of packets per flow when number of flows are large?

↳ Approximate value → compress the data & store approximate

Is this a New Problem?

- **Big data** caused similar problem
 - Billions of records generated per second (e.g., social media logs)
- Datasets too large to fit in memory
- How did big data folks handle this problem?

Data Streaming Problem

- **Data stream:** a sequence $A = \langle a_1, a_2, \dots, a_m \rangle$, where the elements of the sequence (called tokens) are drawn from the universe $[n] = \{1, 2, \dots, n\}$
- **Aim** - compute a function over the stream, eg: median, number of distinct elements, longest increasing sequence, etc.
- **Target Space complexity**
 - Since m and n are “huge,” we want to make s (bits of random access memory) much smaller than these
 - Specifically, we want s to be sublinear in both m and n .



Use Sketch

What are Sketches?

- Data structure that support approximate computing some function of data
- Summarize data into a much smaller space
- Two operations: **update** and **query**
- Common sketches
 - ↪ • Set membership
 - ↪ • Counting

Set Membership Task

$|S| : m$
 $f: U \rightarrow \{0,1\} : m \rightarrow O(m)$

$\uparrow \in U$

- **Problem statement:** Given a set S , does a given query element belong to S ?
- How would you solve this problem in linear space?


space complexity: $O(m)$: ~~not~~ $|S|$

- How to solve this in sub-linear space?

$h: U \rightarrow \{1, \dots, w\}$
 $w \ll |U|$ $S: h(s) = 1$
Bloom Filter \rightarrow False positive


• Query:

k hash functions, 1 boolean vector of $|w|$



k - boolean
 vector
 or k - hash map

$h1: \text{mod } 10$
 $h2: \text{mod } 10$



update (25)

update (21)

Query(22)
 . 3 and 1st bit

Bloom Filter

- Consists of:
 - A vector of n Boolean values, initially all set false
 - k independent hash functions, h_0, h_1, \dots, h_{k-1} , each with range $\{0, 1, \dots, n-1\}$

F	F	F	F	F	F	F	F	F	F
0	1	2	3	4	5	6	7	8	9

$n = 10$

Bloom Filter ^{0'} → # of distinct element in sub-linear space

- Consists of:
 - A vector of n Boolean values, initially all set false
 - k independent hash functions, h_0, h_1, \dots, h_{k-1} , each with range $\{0, 1, \dots, n-1\}$ $\Theta(n)$
- **Update:** For each element s in S , the Boolean value with positions $h_0(s), h_1(s), \dots, h_{k-1}(s)$ are set true
- **Query:** True if $h(s) = T$ for all hash functions

else false

F	F	F	F	F	F	F	F	F	F
0	1	2	3	4	5	6	7	8	9

$n = 10$

Counting

- How to count the number of occurrences of an element s ?

Naive
algorithm! $\log_2 n$ bits

n floors, k



$\rightarrow n \log_2 \log_2 k$

\times

Probabilistic Counting

$\frac{1}{2^x}$

Emulate a counter subject to $X := X + 1$.



Final value of counter was n

$2^n - 1$

Space: $\log_2 \log_2 n$

$\log_k \log_2 n$

$\frac{1}{k^x}$

Count Min Sketch

d hash functions each of length w

update $(h_j(f_i)) \quad \forall j \in \{1, \dots, d\}$

$\oplus, \otimes, \otimes, w-1$

Query (f_i)

$h(f_i) : \min (h_j(f_i)) \quad \forall j \in \{1, \dots, d\}$

→ Heavy-hitters

→ Elephant flows (heavy hitters)

→ Mice flow (not important)

$O(dw)$

1	0	2	0		0
0	1	0			0
0	0	1			0
0	0	0	...		1

Find Top-k flows

Count Min Sketch

- Two-dimensional array counts with width w and depth d
- Each row corresponds to a hash function chosen uniformly at random from a pairwise independent family which map vector entry to $[1...w]$.
- **Update:** For an element s , compute the hash values and increment the corresponding entries in the 2-D array
- **Query:** For an element s , compute the hash value and return the min across all values

Tribe h
→ next - heap

Hash function 1	→	0	0	0	0	0
Hash function 2	→	0	0	0	0	0
Hash function 3	→	0	0	0	0	0

Special Topics: Machine Learning (ML) for Networking

COL867

Holi, 2025

Network Telemetry

Tarun Mangla

Recap: Network Telemetry

- Conventional telemetry is bottom up. High overhead
- Top-down telemetry or query-driven telemetry
- **Challenge**: Can't run arbitrary queries, even on programmable data plane
 - Memory and compute bottleneck
- **Solution (for memory bottleneck)**
 - Use sketches – memory-efficient data structures for summarizing the data
 - Example sketches: bloom filter, probabilistic counting, count min

Implement the following telemetry using sketches

Exercise 2: Count the number of unique sources accessing a web service via port 80

Implement the following telemetry using sketches

Exercise 2: Have you already seen a given unique source s accessing your web service via port 80

Bloom filter

Filter on port 80:

Implement the following telemetry using sketches

Exercise 3: Count the number of unique sources accessing a web service via port 80

Exercise

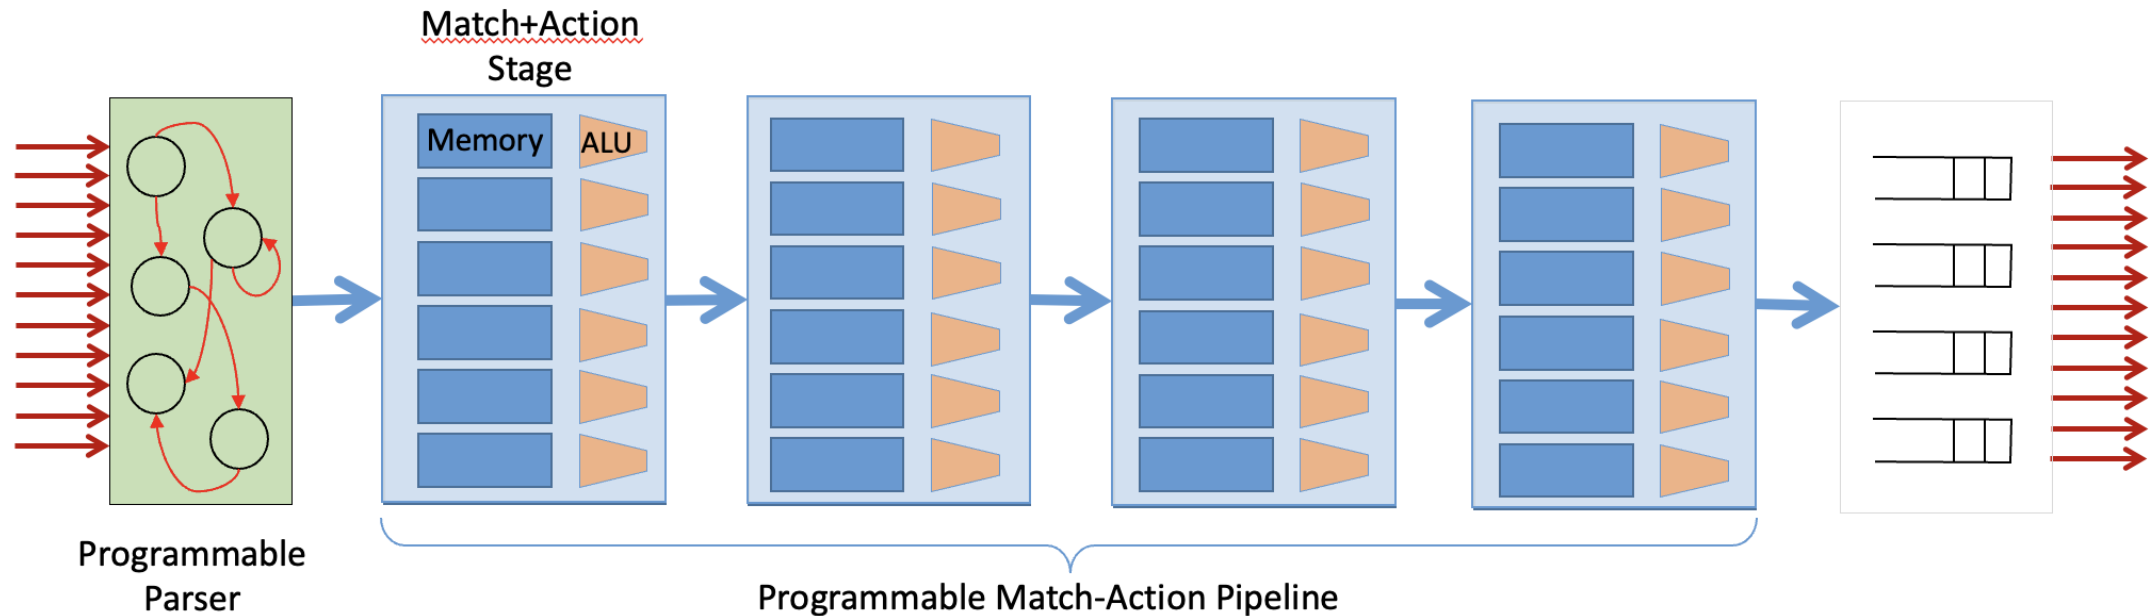
Exercise 4: Identify the heavy hitter flows, i.e., flows that consume more than a fraction T of the link capacity during a time interval

How to implement the sketch in a data plane?

What are the fundamental operations needed for sketch?

- Hashing
- Match
- Probabilistic counter

Programmable Data Plane Architecture



Key functionalities

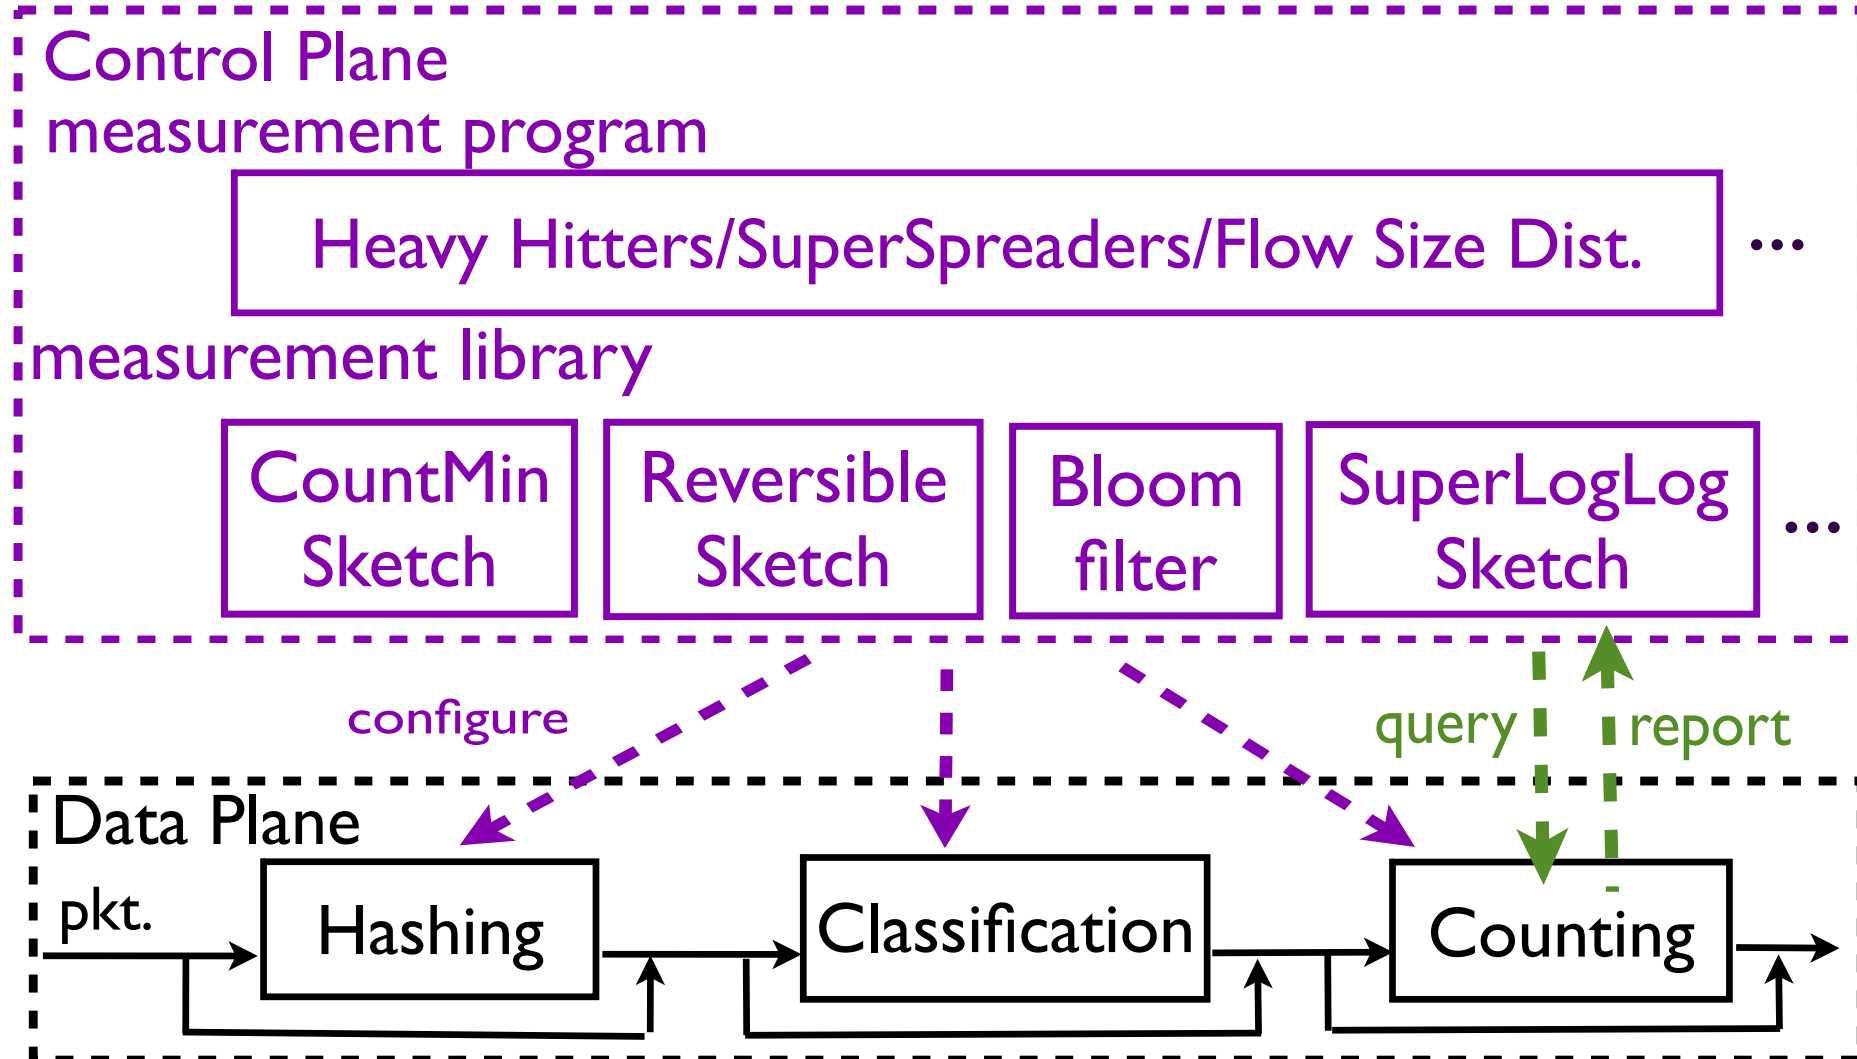
- Programmable parser, extract specific fields of interest
- Matching, to group related packets into a single 'flow'
- Computation using simple ALUs (addition, subtraction, bitwise operations)
- Storage of information across successive packets using a small number of register arrays
- Communication with a software controller

Build on Existing Switch Components

- A few simple hash functions
 - 4-8 three-wise or five-wise independent hash functions
 - Leverage traffic diversity to approx. truly random func.
- A few TCAM entries for classification
 - Match on both packets and hash values
 - Avoid matching on individual micro-flow entries
- Flexible counters in SRAM
 - Logical tables with flexible indexing
 - Access counters by addresses

0	0	2	0	0	0	0	1	1	1	0
0	0	0	2	0	0	0	0	1	1	1
0	3	0	0	0	2	0	0	3	0	1
0	0	1	2	0	0	0	3	0	3	2

OpenSketch Architecture



Discussion

- What about performance metrics like latency?
 - Nice survey paper: [Compact Data Structures for Network Telemetry](#)
- What is the next logical step?